Documentation for **SCAT**, version 3.0.2

Software code and documentation by Matthew Stephens[1]
WWW: `http://github.com/stephens999/scat`

September 28 2004; updated September 15 2015; revised by Mary
Kuhner[2] July 29 2021

[1]email: `mstephens@uchicago.edu`
[2]email: `mkkuhner@uw.edu`

Contents

1 Introduction

The program `SCAT3` (Smoothed and Continuous AssignmenTs) implements a Bayesian statistical method for estimating allele frequencies and assigning samples of unknown (or known) origin across a continuous range of locations, based on genotypes collected at distinct sampling locations. In brief, the idea is to assume that allele frequencies vary smoothly in the study region, so allele frequencies are estimated at any given location using observed genotypes at nearby sampling locations, with data at the nearest sampling locations being given greatest weight. The software can deal with SNP and microsatellite and other multi-allelic loci, in any combination, and missing data are allowed. Details of the method are given in Wasser et al. (2004).

The program is fairly slow: data sets attempting to locate a thousand samples can take 3-4 days to run on a moderately fast machine. We recommend trial runs with fewer samples to get an idea of the time and space requirements.

This document describes the use of this software, which comes free of charge, and with no warranty whatsoever. Updates for the software will be made available via

                    http://github.com/stephens999/scat

Please send bug reports and requests for new features by opening an issue on that page.

1.1 How to cite this program

In publications including results from the use of this program, please *specify the version of the software you used*, and cite Wasser et al. (2004):

Wasser SK, Shedlock AM, Comstock K, Ostrander EA, Mutayoba B, Stephens M (2004) Assigning African elephant DNA to geographic region of origin: applications to the ivory trade. PNAS 101: 14847-14852.

We would love to hear about any such publications (email mkkuhner@uw.edu).

2 Getting started: compiling and running the software

2.1 Compiling

`SCAT3` was previously distributed as executables for MAC OSX and Linux, but currently only for Linux. Users of other operating systems will need to compile from source.

The Linux executable comes in a gzipped tar file. Extract the files from the `.tar.gz` file by typing

                    gunzip scat.linux.X.tar.gz

                    tar -xvf scat.linux.X.tar

at the command line, where `X` is the version number. This will create a new directory called something like `scat.linux.xx`. Change to this directory before running the program.

Source code is available on the github page. To build the program from source, unpack the source code into a directory and change directory into scat-master/src/. Then type `make`.

The Makefile as distributed makes an optimized version of the program for speed. If you need to run the program in a debugger such as `gdb`, turn off optimization and turn on debugging support by locating the following line in the Makefile:

```
CFLAGS = -O3 -std= -DNDEBUG $(INCLUDE)
```

and inserting a hash (#) at the beginning, and locating the following line:

```
#CFLAGS = -g $(INCLUDE)
```

and removing its hash. Then type `make clean` and `make`. The resulting code will be much slower and we recommend this only if you are using a debugger. The debugging version does extensive checking for internal consistency; if the program runs in optimized mode but aborts in debug mode, we would very much appreciate a bug report.

2.2 Running the program

The program can be used in two different ways:

- To estimate allele frequencies at specified locations, based on genotypes observed at a number of sampling locations.

- To "assign" (estimate the location of origin of) individuals based on allele frequencies estimated using reference samples from a number of sampling locations.

In this section we will illustrate allele frequency estimation. See section 3.5 for instructions on using the program for assignment.

To run the program, you must supply two input files: a genotype file, containing the genotype information, and a location file, containing information on the latitude and longitude of sampling locations. To perform continuous assignment you must also supply a file giving the boundaries of the habitat of the individuals being assigned - see section 3.6. Instructions for how to prepare these files are given below. Finally, you must create a directory or folder to hold your results.

Once these files are created you can run the program using

```
./SCAT3 <genotype file> <location file> <outputdir> L
```

where `L` is the number of loci in the genotype file and ¡outputdir¿ is the path to your output directory.

An example genotype file `test.genotype.txt` and an example location file `test.location.txt` are supplied with the software. These have $L = 2$. You can run the program on the test data supplied by placing the `SCAT3` executable and the two input files in the current directory and typing:

```
mkdir results
./SCAT3 test.genotype.txt test.location.txt results 2
```

The program will input the data, perform a number of iterations, and output results in files in the results directory you created. Section 4 describes the output files in more detail.

3   Analysing your own data

To analyse your own data, you must prepare a genotype file and location file in the appropriate format, as described below, and then run the program as above (replacing 2 with the the number of loci in your genotype file). A number of additional options to modify program behavior are described in subsequent sections.

3.1   Genotype file format

The genotype file is supplied by the user to specify genotypes of the individuals to be analysed. Note that `SCAT3` assumes diploid, unphased data (phased data will work but no use is made of the phase information); there must be exactly two input lines per individual sampled.

The file should be a plain text (.txt) file. Its format is similar to that used by the program `STRUCTURE` (see section 7) and can be represented as follows:

```
ID(1) Location(1) Allele(111) Allele(121) Allele(131) ... Allele(1L1)
ID(1) Location(1) Allele(112) Allele(122) Allele(132) ... Allele(1L2)
ID(2) Location(2) Allele(211) Allele(221) Allele(231) ... Allele(2L1)
ID(2) Location(2) Allele(212) Allele(222) Allele(232) ... Allele(2L2)
...
...
ID(N) Location(N) Allele(N11) Allele(N21) Allele(N31) ... Allele(NL1)
ID(N) Location(N) Allele(N12) Allele(N22) Allele(N32) ... Allele(NL2)
```

where

1. `ID(i)` is a string, giving a label for individual i.

2. `Location(i)` is an integer specifying the "Location number" of the sampling location of individual i (the location number for each location is specified in the location file; see below). Individuals of unknown origin should be given a location number of -1.

3. `Allele(il1),Allele(il2)` are integers representing the two alleles in individual i at locus l. **Missing alleles should be represented by -999.** (Programmatically, any negative number will treated as missing data, but we strongly recommend use of -999 for consistency with other programs that use this format.)

See the example genotype file, `test.genotype.txt`, for an illustration.

The `-C` option can be used to skip extraneous columns in this file. It takes an argument of the number of columns to skip; columns will be skipped starting with the first column after the location number.

Some things to note:

1. There should be no commas separating columns, just white space.

2. Individual IDs must not contain any spaces.

3. Any non-negative integer can be used to denote the alleles at a locus. For SNP alleles the obvious choice is 0 and 1, but any two numbers can be used, and it need not be the same two numbers in each column (so recoding SNP data using A=1, C=2, G=3 and T=4 is fine). For microsatellites you can use allele length, or number of repeats, or indeed any set of integers. No use is made of the relative values of these numbers; they are arbitrary labels.

4. The software assumes independence of loci (no linkage disequilibrium), so it is advisable to avoid markers that are very closely linked. Use of closely linked markers will overstate the amount of independent information available and lead to too-tight estimates of confidence. While you can input phased data to `SCAT3` , the fact that you were able to phase it is a strong warning that it may be too tightly linked to work well.

3.2 Location file format

The location file is supplied by the user to specify the latitude and longitude of sampling locations, and also of any other locations for which the user wishes to estimate allele frequencies. The file has one row for each location, and its format can be represented as follows:

```
LocId(1)  LocNo(1)  Latitude(1)  Longitude(1)
LocId(2)  LocNo(2)  Latitude(2)  Longitude(2)
...
LocId(S)  LocNo(S)  Latitude(S)  Longitude(S)
```

where

1. `LocId(s)` is a string, with no internal spaces or tabs, giving a label for location s $(s = 1, \ldots, S)$.

2. `LocNo(s)` is an integer specifying the "Location number" of sampling location s. These location numbers can be any S distinct integers. Typically they will simply increase from 1 to S (so the location number of sampling location s will be s).

3. `Latitude(s) Longitude(s)` are the *decimal* latitude and longitude of location s. Note that N and E are entered as positive numbers, and S and W are entered as negative. So, for example, $47°39'N$ and $122°18'W$ become 47.65 and -122.30.

See the example location file `test.location.txt` for an illustration.

If the -Z (Sectors) option is in effect, an additional column must be added to the location file:

```
LocId(1)  LocNo(1)  SectorNo(1)  Latitude(1)  Longitude(1)
```

3.3 Echo input data to screen (`-E`)

The `-E` option will echo the input genotypes file, location file, and related information to screen at the start of the run. This is useful in verifying that the data are read in correctly. It also prints more verbose progress reports as the run goes on.

3.4 Estimating allele frequencies

To estimate allele frequencies at all locations in the location file, simply run the program as in section 2 above, using

```
./SCAT3 <genotype file> <location file> <outputdir> L
```

3.5 Assigning individuals (`-A, -M`)

To estimate the location of individuals ("assignment") you will need to set either `-A` or `-M` (the "assignment options"). The only difference between these options is that `-A` takes the individuals to be assigned from the main genotype file, and `-M` takes them from a separate file. One or the other may be more convenient for your pipeline.

Both assignment options perform two types of assignments: *smoothing assignment,* which assigns individuals to existing locations, and *continuous assignment,* which assigns them to geographical coordinates that may not coincide with any existing location.

WARNING: If you set neither assignment option, smoothing assignments will still be produced; but no cross-validation will be done, and the certainty of the assignments will be inflated as a result.

Smoothing assignment uses allele frequencies estimated at each location in the location file to compute the probability of each individual's genotype in each location, and outputs the results in the _probs file (see section 4.1). Note that

these probabilities are also computed even if no assignment option is used; the difference is that when an assignment option is invoked these probabilities are computed using leave-one-out cross-validation (that is, the genotypes of individual i are ignored when estimating allele frequencies for assigning individual i).

Continuous assignment uses allele frequency estimates to assign individuals, allowing an individual's location of origin to be anywhere in the organism's range (in particular, the location may be somewhere other than any of the locations in the location file). The organism's range should be specified using a boundary file, as described below. The software outputs a number of estimates for the possible location of each individual, one for each iteration – both burn-in and main iterations – of the MCMC scheme. The results for the burn-in iterations (the first 100 values under the default parameter, since `Nburn=100` by default - see section 5) should generally be discarded. The median latitude and longitude of the remaining samples can be used as a point estimate for the individual's location, and the spread of the points gives an indication of the uncertainty. The results for each individual are saved in a file whose name is the same as the individual ID. (The third column in this file gives the log likelihood of each estimate, which can be used to investigate mixing.)

3.5.1 Assignment from main genotype file (`-A`)

The `-A` option is used when individuals to be assigned are included in the genotype file, on consecutive lines. Individuals of unknown origin should be given a location number of -1 in the genotype file. If an individual is marked for assignment but given a location other than -1, it will participate in allele frequency estimation for that location except for the purposes of its own assignment, where its frequency contribution will be dropped.

To perform assignments for individuals i to j in the input file, simply add `-A i j` after the `SCAT3` command. For example, to perform assignments on the first, second and third individuals in the genotye file use

```
./SCAT3 -A 1 3 <genotype file> <location file> <outputdir> L
```

3.5.2 Assignment from separate file (`-M`)

This option behaves exactly like `-A`, except that rather than being included in the input file, the individuals to be assigned are in a separate file, whose name you specify straight after the `M`, with no space between `M` and the filename: `-Massignmentfile.txt`

The `-M` option does not take arguments; it will always assign all individuals in the assignment file. The format of the assignment file is identical to that of the genotypes file.

3.6 Specifying the organism's range

When performing continuous assignments it is *highly* recommended that you specify the range within which individuals might be found. Without this individuals may be assigned to unrealistic locations (eg land-dwelling organisms assigned to ocean). If only smoothing assignment is needed, this is not necessary as individuals will only be assigned to sampling locations.

Two methods of specifying the organism's range are provided: as a grid indicating all 1 degree squares within the area, or as a polygon bounding the area. The former is recommended, and *must* be used if the SCAT3 data are intended as input to the VORONOI program. (The organism's range must be the same in SCAT3 and VORONOI, and VORONOI cannot read the polygon boundary format.)

Not all organism ranges can be gracefully indicated by the current code: see section 6. In particular, the range cannot include either pole or cross the line opposite the Prime Meridian.

3.6.1 Grid file (-g)

The organism's range can be specified by dividing the area of interest into grid squares of 1 degree latitude and longitude, and listing all grid squares which could contain the organism. An advantage of this method is that it allows a discontinuous habitat, for example an island organism found on several islands but not in the ocean. If the organism has a very large north-south range the distortion of the 1 degree grid squares may be problematic.

To specify the habitat in this way, prepare a file which contains one line per square included in the habitat, giving decimal latitude and decimal longitude with a space between them. The latitude and longitude describe the lower left corner of the grid square, so that an entry:

    -12 17

indicates the grid square whose lower left corner is 12 S and 17 E, and whose upper right corner is 11 S and 18 E.

The grid file is then specified to the program using the -g option followed by the filename of the grid file:

    -g ../savannah_grid.txt

Note that there is a space between the option and the filename.

The program's search area for continuous assignment will be formed by determining the minimum square capable of containing all of the specified grid squares and add a 7 square margin around the edges to avoid edge effects. (This margin size is specified as MARGIN in file scat3.hpp, if you need to change it.)

3.6.2 Boundary file (-B)

The region can be specified as a polygon by entering the latitude and longitude of each consecutive vertex into a "boundary" file (see below), and using the -B option to indicate the name of this file. For example, if the boundary file is

called `eg.boundary.txt` you should run `SCAT3` using

```
./SCAT3 -Beg.boundary.txt ...
```

Note that there is no space between the B and the file name. The boundary file should contain one row for each vertex in the polygon, with two numbers on each row (decimal latitude and longitude, with N and E being positive, S and W being negative) separated by a space. The vertices should be entered in order (in either direction around the polygon), with coordinates of the last vertex being an *exact* repeat of the coordinates of the first vertex.

For example, to specify a square region, from $5°N$ to $3°S$ and $2°W$ to $1°E$ the file could be

```
5.0 -2.0
5.0  1.0
-3.0 1.0
-3.0 -2.0
5.0 -2.0
```

The routine that tests whether a point is inside or outside the polygon is based on the 2-D algorithm helpfully described by Dan Sunday (`http://www.softsurfer.com`).

3.6.3 Pre-loaded Afrian elephant location data

The program contains pre-loaded boundaries for savannah and forest African elephants. These boundaries are not accurate (in particular the savannah boundary contains a good deal of ocean) and we do not recommend their use for any purpose other than replicating previous results.

-d use hard-coded boundaries for savannah African elephants

-D use hard-coded boundaries for forest African elephants

4 Interpreting the output

4.1 Output files

The program produces a number of output files in the specified output directory, whose names are of the form `Output_xxx`. Their contents are as follows:

_freqs Contains estimates (posterior means) of allele frequencies at each location in the location file, for each locus. It also contains, for comparison, the empirical allele frequencies at that location.

_probs Contains estimates of the posterior probability:

log(Pr(genotype data for individual $i|i$ came from location j))

for every i and j. These correspond to assignment to discrete regions, and the common practice is to assign each individual to the location

that maximises this probability. Each row of the file contains the data for a single individual. The first two columns give the individual's id, and the number of loci for which it had genotype data. The next two columns give the location numbers of the true and assigned locations for that individual (the true location will be -1 if not known). Subsequent columns give the estimates of the log probabilities for locations $j = 1, 2, \ldots$. In version 2.1 these probabilities were in hexadecimal; they are now in decimal. In general you will want to use the output from this file *only* if you used an assignment option (`-A` or `-M`). The file is still produced even if you don't use an assignment option, but in that case the results are not based on cross-validation, and assignment results will be inappropriately overconfident.

_params Contains sampled values of model parameters, and log-likelihood of genotype data, for each iteration of the MCMC (parameters are output during both burnin and main iterations, every `Nthin` iterations). There are 5 columns, corresponding to sampled values of $\alpha_0, \alpha_1, \alpha_2, \beta$, and the log-likelihood for the corresponding allele frequency estimates. This file works well with the Tracer utility (see section 7) which can display plots of the parameter values over time. This is helpful in diagnosing poor mixing.

_accept Contains summary of (cumulative) acceptance rates during MCMC runs. Each line contains data for one iteration of the MCMC (after thinning). Acceptance rates on each line are, in order, for α_1, α_2 (if these are updated), X and μ.

_corr Contains estimated and fitted correlations and covariances between each pair of locations. The first two columns give the pair of locations, and the third gives the distance between them in kilometers inferred from their latitude and longitude. The following four columns are empirical and fitted covariance and empirical and fitted correlation. "Empirical" results are estimated from the inferred allele frequences of the two locations. "Fitted" results are estimated using only the distance between the locations and the inferred model parameters (α) which control the rate at which correlation decays with distance. This file may be useful for model checking, specifically spotting pairs of populations that are much more or less correlated than expected based on their geographical distance.

4.2 Individual assignment files

If either assignment option is set, the program will also write one file per sample analyzed, with a filename of the sample name. These files contain results from the continuous assignment algorithm, one result per line, for both burn-in and standard iterations (burn-in iterations should generally be discarded). Each line contains, in order, the decimal latitude, decimal longitude, and log-likelihood

of that sample's assignment for that iteration. The medians of the non-burn-in latitudes and longitudes can be considered point estimates of the sample's origin, or the entire distribution can be considered as an area estimate.

4.3   Verbose output (`-v`)

This option causes the program to write an additional output file containing the contents of various internal arrays including the X arrays, region permutations, counts of observations, and empirical allele frequencies. The name of the output file is given after the `-v`, with a space between them. This is likely useful only to someone trying to modify or debug the code.

5   Search parameters

5.1   Run length

`SCAT3` employs an iterative scheme to perform inference. Parameters that control the number of iterations performed can be added to the input line, as follows:

```
./SCAT3 <genotype file> <location file> <output filename> L Niter Nthin Nburn
```

where `Niter` is the number of sampled iterations of the MCMC scheme to be performed, `Nthin` is the number of steps taken through the Markov chain between samplings (the "thinning interval") and `Nburn` is the number of burn-in iterations.

For example, to set

$$\texttt{Niter=100}, \texttt{Nthin=10}, \texttt{Nburn=100}$$

use

```
./SCAT3 <genotype file> <location file> <output filename> L 100 10 100
```

In fact, the above values are the default values. Each iteration performs `Nthin` steps through the Markov chain. The number of iterations required to obtain accurate answers depends on the complexity and size of the data set; preliminary runs are useful to establish a good run length.

Note that if the VORONOI program will be used for post-processing, it will expect to see 100 burn-in and 100 normal sampled iterations; thus, making the run longer or shorter will require modifiying `Nthin` so that the same number of sampled iterations represented a larger or smaller total number of iterations.

5.2   Setting the seed (`-S`)

The `-S` option can be used to set the seed of the pseudo-random number generator. The seed should be a positive integer, and must follow the `-S` after a space, as in:

```
./SCAT3 -S 3253 ...
```

which will set the seed to be 3253. This option can be used to deliberately duplicate a previous run, or to make sure that you do *not* duplicate a previous run. If the seed is not set, a random number seed based on the computer clock will be used; this can lead to multiple runs getting the same seed if they are started at almost exactly the same time, for example by a batch program.

If you are trying to debug an issue with the program we strongly recommend setting the seed so that you can replicate your run. The execution path of SCAT3 can vary quite a bit from one run to the next if the seed is not the same.

5.3   Isolation by distance parameters

SCAT3 uses a model of isolation by distance for the relationship of different regional populations' allele frequences: there is expected to be more similarity between populations that are closer together. The behavior of this isolation-by-distance model is controlled by three α parameters: α_0 controls the degree to which a regional population varies from the expectation established by other populations, α_1 controls the scaling for the distances, and α_2 controls how quickly correlations between populations drop to 0 with distance. These parameters are estimated by SCAT3 during the run. It also estimates a parameter β which captures the variability of marker loci. For a more detailed explanation, see Wasser et al. (2004).

Two approaches can be taken to these parameters. Initial non-assignment runs can be used to estimate them, as done in Wasser et al. (2004), and then they can be fixed to the estimated values. One can additionally assume that the estimated values can be re-used for other samples from the same species. Alternatively, the parameters can be co-estimated along with the assignments, which probably requires more MCMC steps for accuracy.

5.3.1   Initializing the parameters (-i)

The α and β parameters can be initialized to desired values using -i followed by, in order, values for $(\alpha_0, \alpha_1, \alpha_2, \beta)$ separated by spaces. If good values are known, this might speed up the process of finding better ones.

5.3.2   Estimating and then fixing the parameters (-f, -b)

In Wasser et al. (2004) the program was run with a large number of iterations (Niter=100, Nthin=1000, Nburn=100) and the results in the _param file were used to estimate the α and β parameters. Such runs are best done without assignment, for speed. The parameters were then fixed to their estimated values for the assignment runs. This is done by using -f, followed by a space, and then the values to use for $(\alpha_0, \alpha_1, \alpha_2, \beta)$ (four values, each separated by spaces). For example,

<p align="center">-f 0.43 5300 0.32 2.3</p>

will fix $\alpha_0 = 0.43, \alpha_1 = 5300, \alpha_2 = 0.32, \beta = 2.3$.

Note that when fixing α and β in this way, it is recommended that you first find estimates for α and β from the _params file (e. g. their mean, median or mode after discarding burnin), and then choose the row of estimates in the _params file that is closest (in some sense) to these estimates. This will help ensure that the *combination* of parameter estimates used is somewhat sensible. It is also helpful to do several initial runs, with different values of the seed for the random-number-generator (see -S option below), to check how many iterations are necessary for the values of α and β to be reliably estimated across runs. If there are substantial differences between parameter estimates, or in the range of log-likelihoods achieved, in different runs (output in the _params file) try increasing the lengths of the runs by increasing either the number of iterations (Niter) or the thinning interval (Nthin). The program TRACER (see section 7) can be helpful here. If traces of the parameters appear to move sluggishly, increase Nthin; if there are directional trends visible in the trace over time, increase either Niter or Nthin; if the estimated Effective Sample Size is low, increase Niter.

The program also contains a menu option -b to fix β only, while allowing the α parameters to be estimated. If this is used alone it will fix β to its default of 1.0, as found in scat3.hpp. The -i option can be used to provide a different value.

5.3.3 Co-estimating assignment and parameters

The Center for Conservation Biology's current approach to using SCAT3 is to estimate the α and β parameters during assignment. This was not feasible in 2004 but has become more so, although assignment runs done this way can take several days to complete for around 200-300 individuals. Our current practice is to use Niter=100, Nthin=20, Nburn=100. (If the postprocessor VORONOI is to be used, Niter should always be set to 100.) We run 9 replicates of each SCAT3 run with different random number seeds, both to meet VORONOI requirements and to assess consistency across runs. It is probably better to do 9 medium length runs than 1 very long run, as the variation in starting points can help the program search its entire space more effectively.

5.3.4 Modifying the search in α (-a)

You can also modify the aggressiveness of the search across α values using the following option:

-a sd set the proposal standard deviation for alpha updates to sd (default 0.4). A larger sd will correspond to a more wide-ranging search of alpha values.

5.3.5 Turning off estimation of μ (-m)

The parameter vector μ holds estimates of the population-wide allele frequencies at each locus; the individual sampling locations are considered to vary around these frequencies. Normally they are allowed to vary during the MCMC run.

The `-m` option will instead fix them at values estimated from the initial data, and will also fix β at its starting value. This might possibly be useful if the program is otherwise unacceptably slow.

5.4 Setting error parameters (`-e`)

The assignment analysis machinery allows for the probabilities of an erroneous genotype call at one allele (error term δ) and of amplification of only one allele (error term γ). The defaults are set to the values used in Wasser et al. (2004) of $\delta = 0.05$ and $\gamma = 0$. These values can be changed using `-e delta gamma`.

6 Cautions for use of the software

`SCAT3` carries out an MCMC search. If the search is too short or mixes poorly, the results may be inaccurate and their certainty will be overstated. The program `TRACER` (see section 7) may be useful in diagnosing poor searches. Another diagnostic for too-short runs is getting dramatically different results from multiple runs that differ only in their random number seed.

`SCAT3` can only handle diploid data. Coding haploids as one known haplotype and one missing-data haplotype does not work, since a locus with any missing data is disregarded.

The code which sets the boundaries of the organism's range will not work correctly if the range crosses the line opposite the Prime Meridian or includes either pole. Nothing is done by the code to diagnose such problems. The grid boundary approach may also be problematic if the species has a wide north-south distribution approaching either pole, as the size of the grid squares is assumed to be roughly constant and this will not be the case. If such a species needs to be analyzed, we recommend rescaling the latitudes and longitudes to make them tractable, and then reversing the rescaling before interpreting the data.

If the species occurs as widely separated patches, using a grid file to specify correct species boundaries may cause the continuous assignment algorithm to become stuck on the first solution it finds because other solutions are "too far away" in the patchy map. You can diagnose this by making multiple runs with different random number seeds: if each run places a given individual with great certainty, but different runs place the individual in different patches, this problem is likely occuring. Combining estimates across multiple runs will give a better estimate than any single run. Otherwise, it may be necessary to make the map more continuous even if this requires adding areas where the organism is not found, or to use smoothing rather than continuous assignment.

`SCAT3` assumes that markers are unlinked. Tightly linked markers will cause the accuracy of the results to be overstated. Sequencing individual genes and then using all variable sites within them as markers is likely to cause this problem and cannot be recommended as a way to obtain `SCAT3` input data.

SCAT3 assumes that allele frequencies change smoothly with geographic distance. If there are strong discontinuities in gene flow, it will inappropriately use populations which are nearby but separated by the discontinuity to inform each others' allele frequencies. It may be helpful to separate the discontinuous areas into separate SCAT3 runs, as was done for forest and savannah elephants in Wasser et al. (2004).

There is a hard-coded constant for MAXNALLELE, the maximum number of alleles per locus, near the top of file scat3.hpp. It is set to 120 by default, which is sufficient for microsatellite data and overkill for SNP data. If you have loci with more alleles than this, edit the file and then type "make clean" followed by "make". If you are using SNP data it may be helpful to set this constant to 5 (four bases and unknown data) in order to save memory.

The practice of summarizing SCAT3 results as median latitude and median longitude can be misleading if the individual's inferred distribution is patchy or irregular in shape. Displaying a heatmap of all SCAT3 results for that individual may be preferable, and/or reporting the bounding region containing a given percentage of the estimates.

In our hands, SCAT3 does a poor job localizing individuals which are hybrids between distant populations or species. We recommend the use of EBhybrids (see section 7) to remove putative hybrids. Our preliminary studies suggest that hybrids in the reference data do relatively little harm, but the inferred locations of hybrid individuals are unreliable.

Occasionally the heatmap of an individual may resemble a thin rim around the very edge of the organism's range. In our experience this may mean that the individual is a hybrid, or that it comes from a location which has been incorrectly excluded from the range. In either case, the program finds that the individual is a poor fit for all sampling locations, so moves it as far away from all of them as possible, leading to the thin rim. An alternative explanation is some form of error in the genotype data, such as misordering the columns.

Another unusual outcome is an individual whose continuous assignment is spread across the whole map. If other members of the species can be more narrowly assigned, this individual may be a hybrid, or may have too much missing data. In some cases missing data at one or a few highly informative microsatellite loci can render an individual very difficult to localize.

7 Other software useful in conjunction with SCAT3

These programs are listed as a convenience to the user; we do not guarantee their availability, compatibility, or performance.

- STRUCTURE (http://pritch.bsd.uchicago.edu/) – estimate population admixture and assign individuals to populations.

- EBhybrids (https://github.com/stephenslab/EBhybrids) – postprocess STRUCTURE output to infer hybrid probability for each individual. In our experience, hybrid individuals are difficult for SCAT3 to assign and it

may be best to remove them or at least treat their inferred locations with skepticism.

- **TRACER** (`https://beast.community/tracer`) – display behavior of an MCMC run over time. With minimal processing, the `Output_params` file can be read into `TRACER` to visualize the progress of the run, which is useful in diagnosing poor mixing or too-short runs.

- **VORONOI** (`https://github.com/stephens999/voronoi`) – postprocess `SCAT3` estimates to refine continuous assignment on the assumption that the location-unknown samples may be spatially clustered relative to the reference samples.

8   Options changing the mathematics or search strategy

These options change the behavior of the MCMC sampler. Little information is available about their usefulness, but they are documented here for completeness. In the descriptions below, X refers to the vector of rescaled allele frequencies described in the Online Supplement of Wasser et al. 2004. The `SCAT3` algorithm normally proposes new values of X one by one, with a variance of 0.5, but this can be changed using the options below.

-h sd sets the proposal standard deviation for X updates to `sd` times $\sqrt{1/\alpha_0}$ (default = 0.5). A larger `sd` will correspond to a more wide-ranging search of X values. The higher the value, the further the sampler will attempt to move in one step. Higher values may help in preventing the sampler from becoming stuck (diagnosed by persistent differences between the outcome of different runs on the same data), while lower values may improve acceptance rate (diagnosed by examining the Output_accept file).

-I suppresses permutation of the order in which regions are considered. The program default is to permutate the region order in each run; this causes runs with different random number seeds to process the regions in different orders, increasing the chance that multiple searches will explore different parts of the search space. It is not clear why you would want to turn this behavior off, other than for debugging.

-j update X for an allele jointly, rather than one at a time. No information is currently available on whether this is helpful.

-N includes a "nugget effect" in the covariance matrix. A nugget effect is variation among samples putatively taken at the same point, either because of spatial structure too fine for the sampling scheme to capture, or because of measurement error. No information is currently available on whether this is helpful, but it may be worth considering particularly if your

sample locations are rather broadly defined and you think within-location population structure is likely.

-r use Langevin update for X. This update attempts to improve sampling performance by using the estimated gradient of the local probability density to choose the direction in which to move. The Wikipedia article "Metropolis-adjusted Langevin algorithm" describes this algorithm. No information is currently available on whether this is helpful, but it may be worth trying if mixing is poor.

-R remove all samples from a region when doing assignment of individuals from that region. This is a more aggressive form of cross-validation, but relies on getting enough allele frequency estimates from nearby regions to fill in for the disregarded frequencies for this region, and seems unlikely to work unless the data are quite rich and the number of regions is fairly large. It may be particularly appropriate if the data consist of a large number of regions, each with only a few individuals, so that region-specific empirical allele frequencies rely heavily on the few available samples in that region.

-w use smoothing only towards the mean, with no spatial component (ie set $\alpha_2 = 0$). This asserts that the species has no spatial structure, so that all regions other than the one whose frequencies are currently being calculated contribute equally to that calculation. This is unlikely for real biological populations, though could possibly serve as a comparison point.

-X "cheat" by initialising location close to true location in assignment runs. This can be used when assigning locations to individuals whose location is already known or suspected. The search will converge more rapidly if it starts from a mostly-correct solution. However, if the search is run for too short a time or mixes poorly, it will return its starting location with excessive confidence. We do not recommend the use of this option to generate publishable results. It could be used diagnostically to test whether the sampler is simply stuck at a poor solution, or is actively moving towards a poor solution.

9   Sectors and hybrids (-H)

The use of SCAT3 for hybrid inference has been superseded by the EBhybrids program (see software list), and we strongly recommend use of that program instead. It is documented here only for completeness. Assignment to sectors can be useful under some circumstances.

Sectors are broad geographic areas containing multiple sampling locations. They are most useful when they capture genetic differentiation not captured by the isolation-by-distance method.

SCAT3 can attempt to identify hybrids either between regions or between sectors. Unless the regions are genetically distinct, hybrid identification will be very difficult; hence the use of sectors. SCAT3 can also attempt to assign

locations to identified hybrids, but this is also problematic. It is often better to remove identified hybrids from the analysis.

Sector assignment and/or hybrid identification are turned on with the -H option where the H is followed immediately (no space) by a numeric option. There are three meaningful options:

- -H1 Hybrid analysis with sectors. The program will estimate, for each specimen, the probability that it is solely from each of the sectors, and the probability that it is a hybrid with ancestry from each pair of sectors. (Combinations of more than 2 sectors are not considered.)

- -H2 Normal analysis with sectors. This estimates the sector membership of each elephant, as in the normal region-based assignment but using sectors instead of regions.

- -H11 Hybrid analysis with regions. The program will estimate hybrid probabilities as in H1, but using regions instead of sectors. This is unlikely to work unless the regions are genetically distinct.

For -H1 and -H2 you will need to specify the sector of each sampling location in the location file, and use the -Z option above.

Sector names must be 0,1,...NSUBREGIONS-1 where NSUBREGIONS is a constant in the file scat3.hpp, defaulting to 6. The program *may* work correctly with fewer than NSUBREGIONS sectors, except that rows and columns corresponding to those sectors will be meaningless; it is likely to crash with more than NSUBREGIONS sectors. To adjust this number, edit scat3.hpp and recompile.

The results of this analysis are written to a file `Output_hybrid` in the output director. This file contains one row per individual.

For -H2 there is one column per sector, and the (i, j) entry is

$$\log(\Pr(\text{ind } i\text{'s genotype data} \mid i \text{ comes from } j)).$$

In the case of the "hybrid" options (-H1 and -H11) the columns are

$$\log(\Pr(\text{ind } i\text{'s genotype data} \mid i \text{ has parents from } (j_1, j_2))),$$

for $(j_1, j_2) = (0, 0), (0, 1), \ldots, (0, 5), (1, 0), (1, 1), \ldots, (1, 5), \ldots, (5, 5)$.

The maximum entry in a row can be taken as an estimate of the individual's origin. For the hybrid cases, the overall probability that an individual is a hybrid can be taken as the sum of all entries in the row except those indicating both parents from the same region or sector.

10   Changes to the software

I presume that the changes listed in the documentation for version 2.1 (September 8, 2004) were from 2.0 to 2.1; this is not entirely clear in the documentation.

10.1 Changes between 2.0 and 2.1

According to the previous maintainer, "A number of small changes were made to the MCMC update scheme. Most of these changes were made to try to improve mixing for general datasets (rather than the specific one we analysed)."

A bug in computation of the likelihood for sectors was fixed.

At some point a change was introduced (probably accidentally) which caused the log probabilities in file `Output_probs` to be printed in hexadecimal; it is not known if this behavior was already in 2.0 or was introduced in 2.1.

10.2 Changes between 2.1 and 3.0

These changes were made by Mary Kuhner and Jon Yamato, with permission of the original author.

10.2.1 Corrections

An initialization bug was fixed. This only disrupted the first burnin step, and likely had little impact on correctness; however, runs after this fix will not be numerically identical to those before, as the MCMC search will take a randomly different path.

The probabilities in file Output_probs were changed (back?) to decimal; this may disrupt scripts written for the previous version.

10.2.2 Improvements

Specification of species boundaries via the grid file was added. This functionality is in an additional source file `readboundary.cpp` with function templates in `readboundary.hpp`: the code is shared with VORONOI so that they will behave exactly the same in this regard.

The hard-coded limits on number of regions and number of loci were removed; the program can now handle any reasonable number of regions and loci and does not need recompilation to do so. Unfortunately, there is still a hard-coded limit on number of alleles per locus, set to 120 in the distribution version; if you have any loci with more than 120 alleles you will need to modify this and recompile. (If you have fewer alleles per locus, and suspect you are running out of space, setting this lower may help.)

Printing of input data to screen at the start of a run is now turned off by default, and option `-E` was added to turn it back on if desired. Progress reports during a run were greatly reduced; `-E` will restore those as well. These changes are meant to reduce the chance that an important message will be lost in reams of routine output.

Increased error checking of input options and data was added. Specifically, each individual is now verified to have exactly two haplotypes, and the number of loci must be the same for all haplotypes of all individuals. Runtime options which take numbers as input now check if they are legal; options which read from files check if the file exists.

Version number was added. The Makefile was revised so that "make" makes the `SCAT3` executable. Function templates were moved to a file `scat3.hpp`. The documentation was extensively rewritten.

Internal variable handling was extensively revised to improve maintainability (mainly changes from C arrays to C++ std::vector). This should not impact correctness but may cause the MCMC search to take a different path due to rounding differences.

10.2.3 Removals

Option `-z` which caused certain individuals to be deleted from the data was removed; it had been added to test a specific hypothesis on a specific data set and was not suitable for distribution. Options which set parameters no longer used in the code were also removed: `-c, -F, -T`. An option `-E` which claimed to allow resetting of the upper limits on loci and alleles (but actually did nothing) was removed and its letter repurposed.

11 Acknowledgements

The software makes use of the LAPACK linear algebra routines for finding the Cholesky decomposition of a matrix, and a version of the `wn_PnPoly()` algorithm by Dan Sunday

(`http://www.softsurfer.com/Archive/algorithm_0103/algorithm_0103.htm`)

for finding which areas are within an arbitrary polygon.